Microservices and FaaS for Offensive Security

Ryan Baxendale

\$ whoami

Ryan Baxendale

Penetration Tester
Centurion Information Security Pte Ltd - www.centurioninfosec.sg
Singapore

twitter.com/ryancancomputer github.com/ryanbaxendale linkedin.com/in/ryanbaxendale

AWS Lambda

Jan 2015 - AWS Lambda Preview Open to all AWS Customers

Upload your code and trigger it to run Scales quickly 1,000,000 code runs for free every month

Only supported node.js

lambdash: AWS Lambda Shell Hack By Eric Hammond https://github.com/alestic/lambdash Run shell commands using node.js



Pay for only the compute time you consume.



Servers are dead...

"Serverless"



No Servers to Manage

AWS Lambda automatically runs your code without requiring you to provision or manage servers. Just write the code and upload it to Lambda.

Run Code, Not Servers - Serverless Computing

(Ad) aws.amazon.com/lambda ▼

Use AWS Lambda To Run Code Without Managing Servers. Get Started Today!

Continuous Scaling · Subsecond Metering · Real-Time Data Processing · No Servers to Manage

Product Details FAQs

One-Year Free Account Getting Started Guide



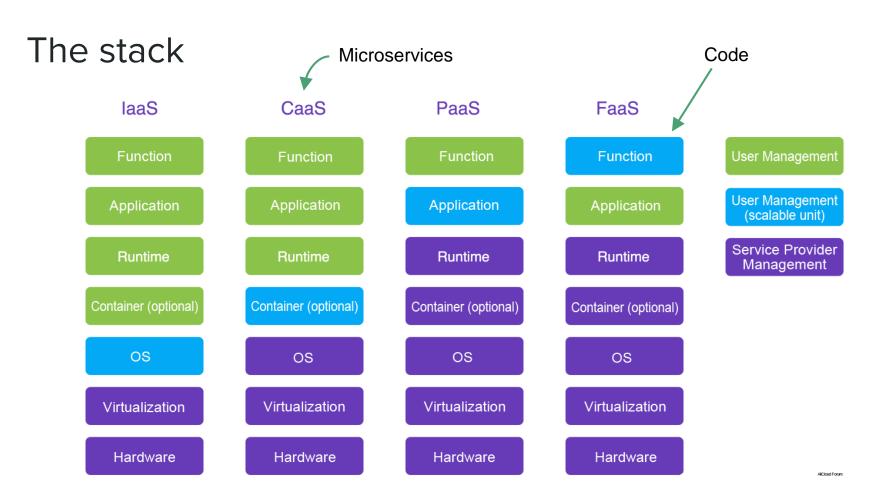
Serverless Economics

Cloud Functions are ephemeral, spinning up on-demand and back down in response to events in the environment. Pay only while your function is executing, metered to the nearest 100 milliseconds, and pay nothing after your function finishes.



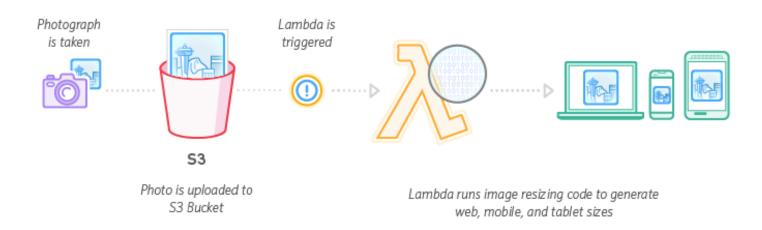
Just Add Code

Run in a fully-managed, serverless environment where Google handles infrastructure, operating systems, and runtime environments completely on your behalf. Each Cloud Function runs in its own isolated secure execution context, scales automatically, and has a lifecycle independent from other functions.



Real-time File Processing

Example: Image Thumbnail Creation



Scale

https://github.com/airbnb/streamalert

StreamAlert is a serverless, realtime data analysis framework which empowers you to ingest, analyze, and alert on data from any environment, using datasources and alerting logic you define.

https://github.com/0x4D31/honeyLambda

honeyλ - a simple serverless application designed to create and monitor URL {honey}tokens, on top of AWS Lambda and Amazon API Gateway

https://github.com/goadapp/goad

Goad is an AWS Lambda powered, highly distributed, load testing tool built in Go

https://github.com/davbo/lambda-csp-report-uri
Simple python application which runs on AWS
Lambda and writes **CSP** reports into S3 for later
processing

https://github.com/therefromhere/csp_lambda

AWS Lambda function to store Content Security

Policy reports in ElasticSearch







Goad

Automate

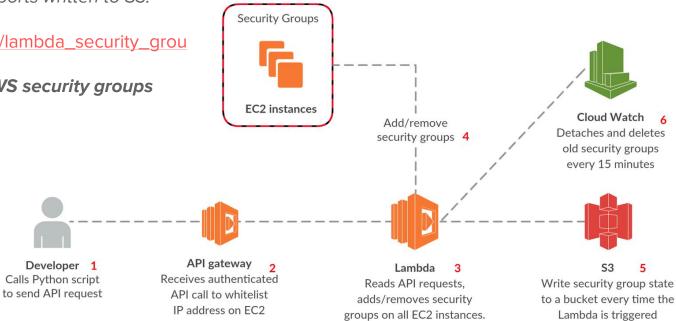
https://github.com/marekq/aws-lambda-firewall Create temporary security groups on your EC2 instances through a simple API call. In addition, audit your security groups easily by the use of automated reports written to S3.

https://github.com/ilijamt/lambda_security_group_manager

Auto managing your AWS security groups with I ambda

https://github.com/johnmccuk/cloudflare-ip-security-group-update

Lambda function to retrieve **Cloudflare's IP address** list and update the specified security group



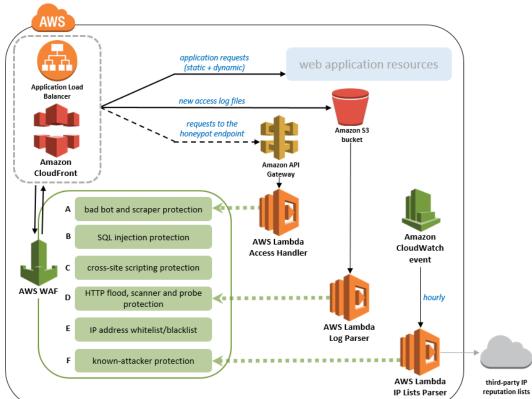
AWS WAF Automation

https://aws.amazon.com/answers/ security/aws-waf-securityautomations/

Parse application logs and trigger WAF rules

Honeypot Log parsing (db scraping) Use third party IP reputation lists





Hello World from the Serverless cloud

■ Google Cloud Platform

Serverless Applications on Google's Infrastructure

Cloud computing has made possible fully serverless models of computing where logic can be spun up on-demand in response to events originating from anywhere. Construct applications from bite-sized business logic billed to the nearest 100 milliseconds, only while your code is running. Serve users from zero to planet-scale, all without managing any infrastructure.



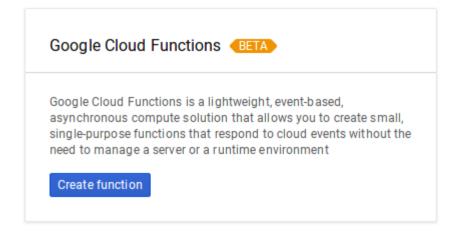
Hello Serverless World

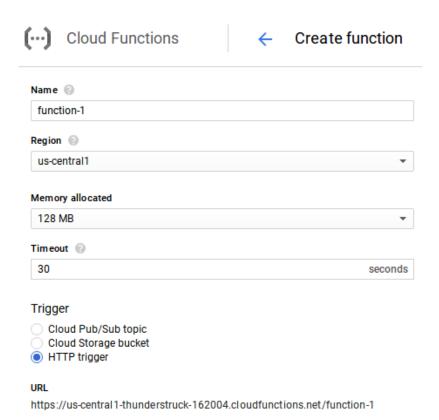
Hello World on Google Cloud Functions (1/4)

Support for Node.js

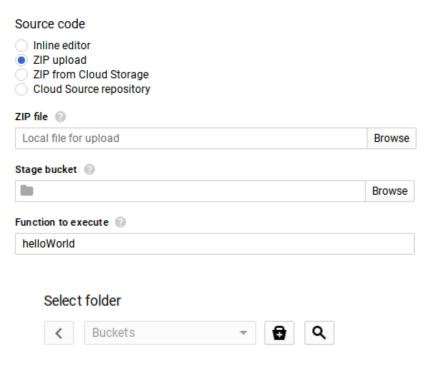
Run Python scripts from Node.js with simple (but efficient) inter-process communication through stdio

https://github.com/extrabacon/python-shell





Hello World on Google Cloud Functions (2/4)

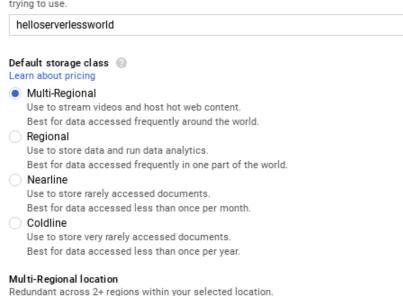


Name 🙆

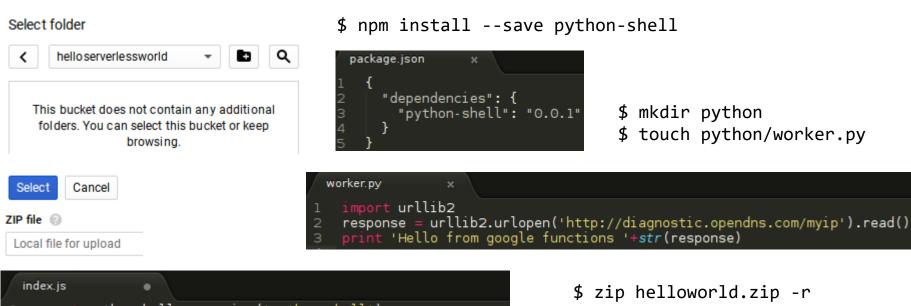
Asia

Cancel

Must be unique across Cloud Storage. Privacy: Do not include sensitive information in your bucket name. Others can discover your bucket name if it matches a name they're trying to use.



Hello World on Google Cloud Functions (3/4)

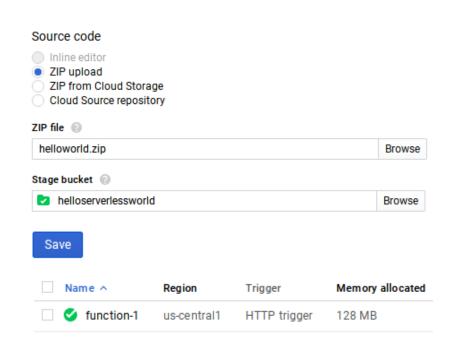


```
index.js

const pythonshell = require('python-shell');
exports.helloWorld = function helloWorld(req, res) {
    pythonshell.run('worker.py', function (err, results) {
    if (err) throw err;
    res.status(200).send(results);
    });
};
```

```
$ zip helloworld.zip -r
python/ index.js
package.json
node_modules/
```

Hello World on Google Cloud Functions (4/4)





IP address is 2600:1900:2000:1c:400::10

IBM Bluemix OpenWhisk

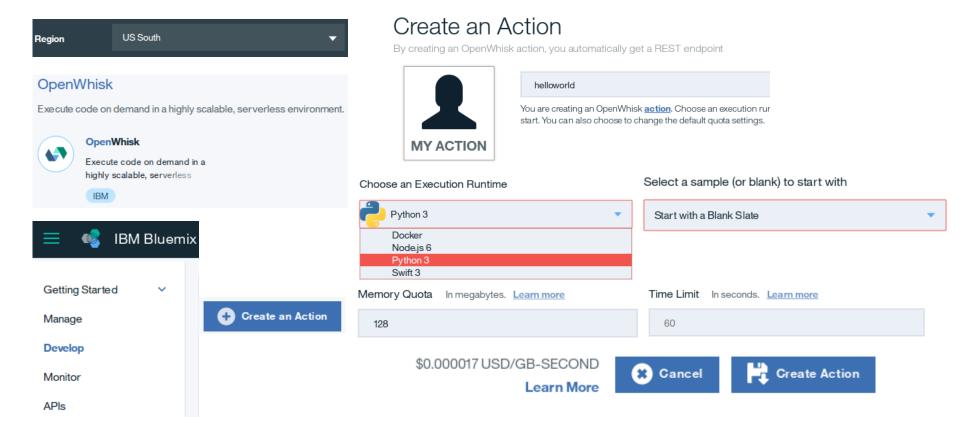
Execute code on demand in a highly scalable serverless environment

Get started free

- Focus on your essential event-driven logic, not on maintaining servers
- Integrate with a catalog of services
- Pay for actual usage rather than projected peaks

Hello Serverless World

Hello World on IBM OpenWhisk (1/2)



Hello World on IBM OpenWhisk (2/2)

Python3

helloworld

```
1 import sys
2 from urllib.request import urlopen
3 * def main(dict):
4    my_ip = urlopen('http://diagnostic.opendns.com/myip').read()
5    print("Hello from azure functions "+str(my_ip))
6    return {'ip': str(my_ip)}
```





```
Invocation Console
 Invoked
 helloworld
                  "ip": "b'169.46.101.173"
 Invoked at
 4:17:02 PM
 Completed in
 903ms
 Billed for
 100ms
```

2017-06-24T08:17:03.26416996Z stdout: Hello from azure functions b'169.46.101.173'

IP address is 169.46.101.173

Welcome!

Before starting we need to verify you are a human



Hello Serverless

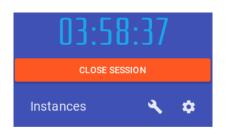
Sessions and all their instances are deleted after 03:59:55 hours.

Hello World on Play with Docker

Hosted: http://www.play-with-docker.com/

Build your own: https://github.com/alexellis/faas

A serverless framework for Docker



+ ADD NEW INSTANCE

\$ curl ifconfig.co 34.206.199.2 \$ python -V Python 2.7.13

\$ dig +short -x 34.206.199.2 ec2-34-206-199-2.compute-1.amazonaws.com.

- +Anonymous (no account)
- -time limited
- -captcha



Cost

Simple Hello World function to get current IP http://serverlesscalc.com/

Serverless Cost Calculator

(beta

Calculating cost for AWS Lambda, Azure Functions, Google Cloud Functions, and IBM OpenWhisk

10000000	Number of Executions	
300	Estimated Execution Til	me
128MB	(ms)	
● True ○ False	Memory Size	
○ True ● False	☑ Include Free-Tier	
	HTTP Requests	

AWS: "1M free requests per month and 400,000 GB-seconds of compute time per month"

128 MB = 3,200,000 free seconds per month Then \$0.000000208 per 100ms

300ms - 10 million executions for \$1.80

Vendor	Request Cost	Compute Cost	Total
AWS Lambda	\$1.80	\$0.00	\$1.80
Azure Functions	\$1.80	\$0.00	\$1.80
Google Cloud Functions	\$3.20	\$4.00	\$7.20
IBM OpenWhisk	\$0.00	\$0.00	\$0.00

FaaS support by region

AWS

- 1. US East (N. Virginia)
- 2. US East (Ohio)
- 3. US West (N. California)
- 4. US West (Oregon)
- 5. Canada (Central)
- 6. EU (Ireland)
- 7. EU (Frankfurt)
- 8. EU (London)
- 9. Asia Pacific (Singapore)
- 10. Asia Pacific (Sydney)
- 11. Asia Pacific (Seoul)
- 12. Asia Pacific (Tokyo)
- 13. Asia Pacific (Mumbai)
- 14. South America (São Paulo)

Azure

- 1. East US
- 2. East US 2
- 3. West US
- 4. West US 2
- 5. South Central US
- 6. North Central US
- 7. Central US
- 8. Canada Central
- 9. Canada East
- 10. North Europe
- 11. West Europe
- 12. UK West
- 13. UK South
- 14. Southeast Asia
- 15. East Asia
- 16. Japan West

<u>Azure</u>

- 17. Japan East
- 18. Brazil South
- 19. Australia East
- 20. Australia Southeast
- 22. Central India
- 23. South India

IBM

1. US South

Google

1. IOWA (us-central1)

Play with Docker

1. (AWS)

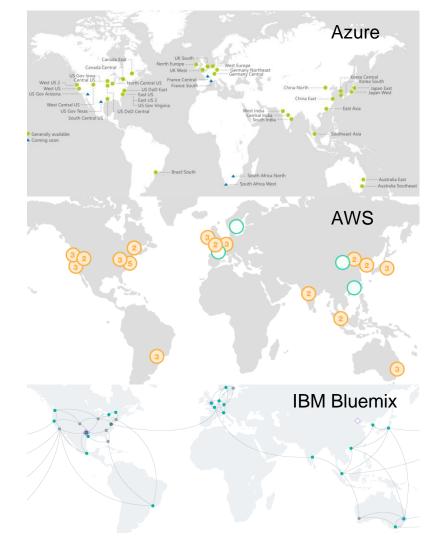
Overview

	Google	IBM	AWS	Azure
Regions	1	1	<u>14</u>	23
Language	Node.js (Python)	Docker Node.js 6 Python 3 Swift 3	Edge Node.js 4.3 Node.js 4.3 Node.js 6.10 Python 2.7 Python 3.6	Bash, Batch C#, F# JavaScript Php, PowerShell Python, TypeScript
OS (Python)	Linux Debian 8.8	Linux Ubuntu 14.04.1	Linux 4.4.51-40.60.amzn1.x86_64	Windows Server 2012
Network	IPv6	IPv4	IPv4	IPv4

Advantages

- 1. Low cost ("free")
 - a. Sign up credit
- 2. Unspecified source IP addresses
 - a. Possibly low attribution
- 3. Global data centers
 - a. China









Project Thunderstruck

Finding use cases for FaaS in offensive security

Project Thunderstruck

Finding use cases for FaaS in offensive security

Explore different cloud service providers

Try to get *supercomputer* resources without paying supercomputer prices

BSidesLV 2017

1. Searching in IPv6

Searching in IPv6 NOERROR but keep looking

Past work

You can -j REJECT but you can not hide: Global scanning of the IPv6 Internet
- Tobias Fiebig

December 2016 33C3: 33rd Chaos Communication Congress, Hamburg, Presenting first steps towards scanning IPv6 globally.

https://media.ccc.de/v/33c3-8061-you_can_j_reject_but_you_can_not_hide_global_scannin g_of_the_ipv6_internet

https://www.cs.ucsb.edu/~vigna/publications/2017_PAM_CollectingIPv6.pdf

https://gitlab.inet.tu-berlin.de/ptr6scan/toolchain

Finding v6 hosts by efficiently mapping ip6.arpa

- Peter van Dijk

March 2012

http://7bits.nl/blog/posts/finding-v6-hosts-by-efficiently-mapping-ip6-arpa

https://tools.ietf.org/html/rfc7707

Standard-Compliant Nameservers (DNS)

Something From Nothing (There): Collecting Global IPv6 Datasets From DNS

Tobias Fiebig, Kevin Borgolte, Shuang Hao, Christopher Kruegel, Giovanni Vigna

PAM (Passiv and Active Measurement) Conference 2017, Sydney

NOERROR

Keep looking

NXDOMAIN

Give up there's nothing here

"Complementary to prior approaches, van Dijk enumerates IPv6 reverse records... When correctly implementing **RFC1034**, as clarified in **RFC8020**, the Name Error response code (**NXDOMAIN** in practice) has the semantic of there is nothing here or anywhere thereunder in the name tree."

"Combined with the IPv6 PTR DNS tree, where each subzone has 16 (0-f, one for each IPv6 nibble) children up to a depth of 32 levels, provides the possibility to exploit standard-compliant nameservers to enumerate the zone"

Let's dig over here

```
$ dig
1.0.0.0.0.0.0.0.0.0.2.0.1.0.0.2.ip6.arpa.
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status:
NOERROR, id: 18347
;; flags: qr rd ra; QUERY: 1, ANSWER: 0,
AUTHORITY: 0, ADDITIONAL: 1</pre>
```

There are records to be found within:

1.0.0.0.0.0.0.0.0.0.2.0.1.0.0.2.ip6.arpa.

2001:0200:0000:0001:...

```
$ dig
5.0.0.0.0.0.0.0.0.0.2.0.1.0.0.2.ip6.arpa.
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status:
NXDOMAIN, id: 58924
;; flags: qr rd ra; QUERY: 1, ANSWER: 0,
AUTHORITY: 0, ADDITIONAL: 1</pre>
```

Nothing to be found within:

5.0.0.0.0.0.0.0.0.2.0.1.0.0.2.ip6.arpa.

2001:0200:0000:0005:...

Super Computer

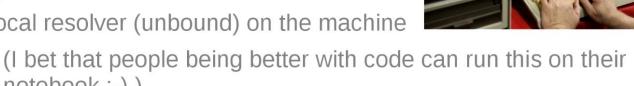
Tobias Fiebig, University computer used for Machine Learning

My Measurement Setup

- 80 Thread HP DL580g7
- 512GB Ram
- 2TB Fast storage
- 1gE Network to DFN

notebook ;-))

Local resolver (unbound) on the machine





Results and Limitations

1st Attempt

Breadth first search 1 week 70M+ records

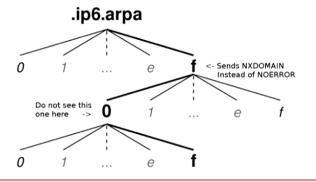
Received complaint email from large ISP Dynamic reverse zones

2nd Attempt

Detect and ignore auto-gen zones 3 days (65.53 hours) 1.6M records

Non RFC8020 Compliance

 Servers that sent NXDOMAIN instead of NOERROR for nonexisting nodes with children:



Not all DNS servers reply with NOERROR

Seeding

Non RFC8020 Compliance

- Solution: Seeding, using an aggregated BGP view
- Algorithm:
 - In: 2001:1:2:3::/64 as seed from BGP
 - Out:

```
3.0.0.0.2.0.0.0.1.0.0.0.1.0.0.2.ip6.arpa.
2.0.0.0.1.0.0.0.1.0.0.2.ip6.arpa.
1.0.0.0.1.0.0.2.ip6.arpa.
1.0.0.2.ip6.arpa.
```

- Other possible sources:
 - Czyz et al.'s PTRv4 -> AAAA lookups
 - All v6 Datasets you can get your hands on!





Scale to Fail

3rd Attempt

70.5 hours80 threads5.3M records

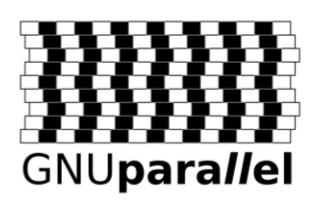
4th Attempt

22.2 hours 400 threads 2.2M records <---:'(

"Lesson: use more than one IP address" "Server ran out of sockets"

Good use case for FaaS?

Bash script curl/wget routing data BGPdump sort --parallel | uniq parallel cook_down.py parallel extract_terminals.py



The Plan

Question to Tobias Fiebig during 33c3:

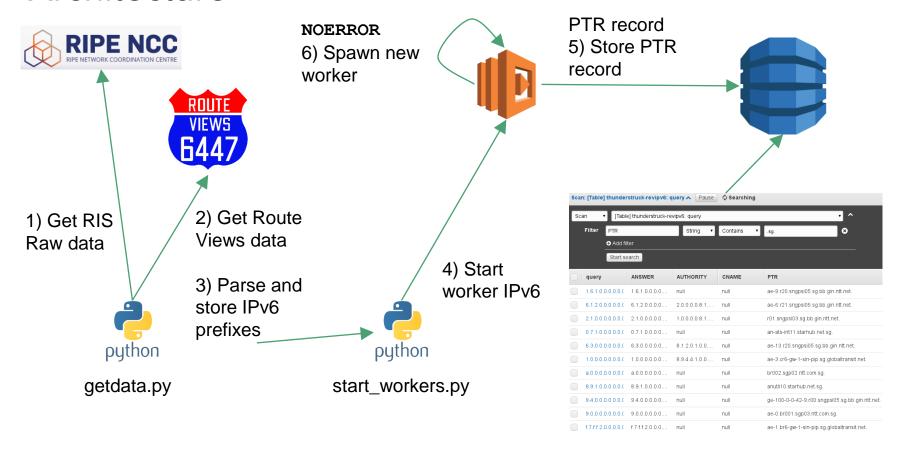
Q: "Where is the dataset?"

A: "It's actually stored in a distributed manner on a lot of DNS servers"



- Download public routing information (BGP) for advertised IPv6 networks
- 2. Parse and create a list of starting IPv6 addresses
- 3. Run worker to enumerate each starting address
- 4. Write any discovered PTR records to DynamoDB
- 5. Look for interesting stuff in DynamoDB web interface

Architecture



Start the workers!

\$./getdata.py

\$./start_workers.py

Need to spawn 55548.0 workers to lookup

55548 with 1 ipv6 address per worker

02:59:47.796904 - starting revipv6

02:59:47.865480 - starting workers

03:05:36.023268 - done starting workers

finished starting workers in 0:05:48.157793

Provisioned read capacity units

10 (Auto Scaling Disabled)

Provisioned write capacity units

100 (Auto Scaling Disabled)

Last decrease time

me -

Last increase time

July 14, 2017 at 8:46:35 AM UTC+8

Storage size (in bytes)

114.82 MB

Item count

250,679

250k+ records found



That was fast, but where are all the records?

Learning to scale

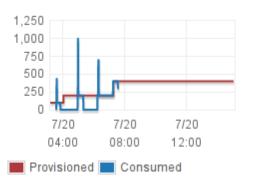
CloudWatch logs:

An error occurred (ProvisionedThroughputExceededExcep tion) when calling the PutItem

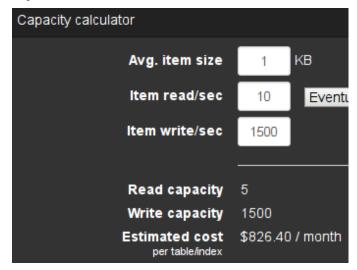
operation (reached max retries: 9): The level of configured provisioned

throughput for the table was

exceeded. Write capacity (Units/Second - 1 min avg)



Too many concurrent attempts to write to DynamoDB



Need to scale the DynamoDB storage backend USD\$ 826/month!

Improvements

Remove the recursive function call

Scales too quickly for too little work

Create depth search limit

 Only look 4, 8, 16, 24, 32 addresses deep from each starting point

Start workers slowly

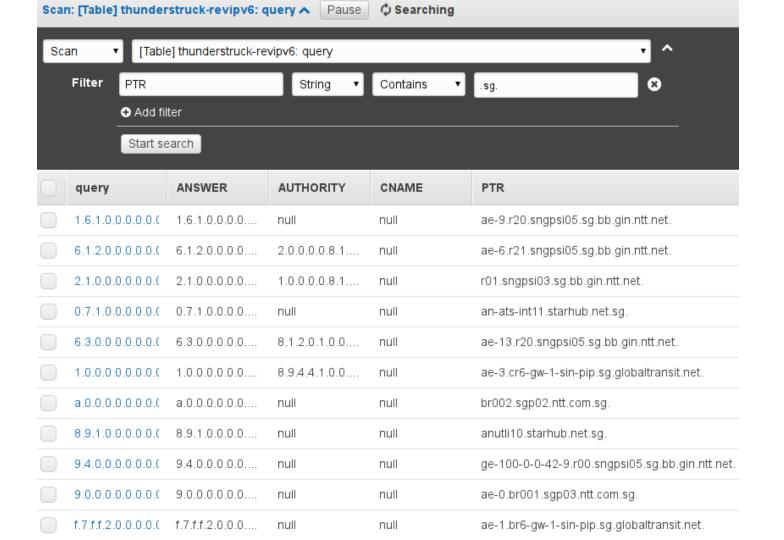
 Don't try to look up all possible reverse DNS entries for every IPv6 address in under 5 minutes

Scale up DynamoDB when needed Scale down when idle

DynamoDB web interface is slow and clunky at search ("scan")



Let's look at some results



KVM

Intelligent Platform Management Interface (IPMI)

ipmi.k7-a14-lt1.kvm.serveriai.lt

ash-ipmi.bootc.eu

esx3-ipmi.sgef.hu

ipmi.origo.nmugroup.se

meter5-test-ipmi.meter.cesnet.cz

dfw1-ipmi.psg.com

vm0-ipmi.iad.rg.net

hippo-ipmi.nist.gov

monitoring-ipmi.dataix.ru

Integrated Lights-Out (iLO)

ilo.app9.mia2.afilias-nst.info

vrrp.ilo.cpr1.lax1.afilias-nst.info

vm11-ilo-gs1-gcn.tgsa.co.uk

tur-**esxfw1-ilo**.massey.ac.nz

OOB

gate-2-oob.belval.restena.lu

oob.core.nrc.ca

interface-3--hoc1-m15-01-hnd.oob.sfdc.net

oob-dc2.ksrv.fr

xrc1-oob.1d12.xs4.mendix.net

internal-240.**fw-oob**.workbone.pironet-ndh.com

eth0.oob-pix.egh.uk.mdnx.com

Integrated Dell Remote Access Controller (iDRAC)

insphl-idrac.spawar.navy.mil

djs-sandiego-proto-idrac.mgmt.dren.mil

idrac-vmhost6.copernica.nl

perez-drac.lusem.edu

drac1.softlab.ece.ntua.gr

pc19-drac.otemachi.wide.ad.jp

Gov

<u>.mil</u>

djs-sandiego-proto-**idrac**.mgmt.dren.mil

fw.sd.spawar.navy.mil

sentinel.tni-au.mil.id blackbox.tni-au.mil.id incognito.tni-au.mil.id

www-largedata-proj.cmf.nrl.navy.mil

remote.mda.mil stafflan.svcs.dren.mil .gov

sec1sgw1.internet.gov.sa gcca-new-transit-core-gcc.fnal.gov

VRRPv6.s3.bnl.gov

ir2-n1--vlan501.**hsrp**.lbl.gov

bluecoat-sg-cadev-3.ca.sandia.gov

Infrastructure

EQIX-EMA-SV-IPV6.softbanktelecom.com

EQIX-EMA-NY-IPV6.bankofnymellon.com

EQIX-EMA-NY-IPV6.centrilogic.com

EQIX-EMA-LA-IPV6.cloudflare.com

EQIX-EMA-**DA**-IPV6.oraclecorporation.com

EQIX-EMA-SV-IPV6.evernote.com

EQIX-EMA-NY-IPV6-HSRP-gateway.ny.equinix.com

EQIX-V6-MA-dc.box.com

EQIX-MA-CH-IPV6.twitter.com

EQIX-MA-SV-IPV6.vmware.com

EQIX-MA-DA2-IPV6.NETFLIX.com

lo0.gw102.me1.ap.equinix.com

xe-0-0-0.gw603.hk3.ap.equinix.com

xe-0-0-0.gw101.ty4.ap.equinix.com

xe-0-1-0.gw401.**sg1**.ap.equinix.com

loop0.gw01.stjh.nf.aliant.net

loop0.cmts01-cres.mql.ncable.net.au

loop0.ubr10k.omnitel.biz

loop0.core1.camphill.dcsi.net.au

loop0.pling.v6.ngdc.net

loop0.23w.ba07.sydn.ns.bellaliant.net

loop0.hkg-mgi-mgmt-1.megaport.com

loop0.sin-sg2-mgmt-1.megaport.com

backbone-v6-r23.ictvalleumbra.it

Quantum Computing

NASA

dwave-**fw**.nas.nasa.gov dwave-**pp2**.nas.nasa.gov dwave-**qubist**.nas.nasa.gov dwave-**monitor**.nas.nasa.gov dwave-**qc**.nas.nasa.gov

heliophysicsdata.sci.gsfc.nasa.gov rtr-s28-40g-v6.sci.gsfc.nasa.gov oceanpost.sci.gsfc.nasa.gov voyager.sci.gsfc.nasa.gov

Summary

Received zero abuse emails

Lessons learned:

Recursive functions that call themself quickly multiple beyond control

Limit recursive calls

Always make sure your backend data storage will scale

 Elasticsearch is expensive to scale and likely can't handle the workload

Going further

More data sources for seeding https://github.com/trustedsec/hardcidr

Public datasets, Certificate transparency, Forward DNS dataset, DNS, scans.io https://github.com/appsecco/bugcrowd-levelup-subdomain-enumeration

Create workflow with DynamoDB Triggers to run another Lambda function (portscan, find matching ipv4, etc)

AWS Lambda free tier is per region

 Reduce cost even further by allocating workload to regions with free compute

Further work



Interesting talks

33c3 2016

Gone in 60 Milliseconds
Intrusion and Exfiltration in Server-less
Architectures
Rich Jones

Blackhat US 2016
Account Jumping Post Infection
Persistency & Lateral Movement In AWS
Dan Amiga & Dor Knafo

Blackhat US 2017

Hacking Serverless Runtimes: Profiling AWS Lambda Azure Functions and more

Wednesday, July 26 - 1:30pm - 2:20pm Andrew Krug & Graham Jones

DEF CON 25

Starting the Avalanche: Application DoS In Microservice Architectures

Friday, July 28 - 13:00 in Track 3 Scott Behrens, Jeremy Heffner

DEF CON 25

Microservices and FaaS for Offensive Security

Saturday July 29 - 11:00 in 101 Track

- 1. DDoS without Servers
- 2. SMS OTP Brute Force

Going further

AWS Lambda - High mem: 1536 MB 266,667/seconds/month free

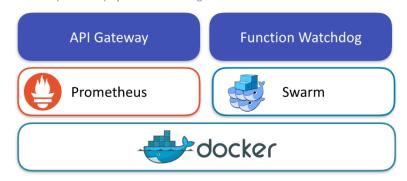
Aliyun / Alibaba Cloud - China Need to register with +86 mobile number

IBM OpenWhisk Docker



FaaS Stack

FaaS is an open-source project written in Golang and licensed under the MIT license.



Build your own FaaS infrastructure

https://github.com/alexellis/faas

UI portal

Setup with one script

Any process that can run in Docker can be a serverless function

Prometheus metrics and logging

Auto-scales as demand increases



github.com/ryanbaxendale/thunderstruck-demo